

Yubo – Cmake, Docker and KAFKA

Compilation without Docker

Here is the file structure for YLibrary :

```
>> cd YLibrary; ll
- Test/src
- YLibrary/include/y/Book
  YLibrary/include/y/Config
  YLibrary/include/y/Event
  YLibrary/include/y/Hitter
- YLibrary/include/y/Oms
  YLibrary/src/Book
  YLibrary/src/Config
  YLibrary/src/Event
  YLibrary/src/Hitter
  YLibrary/src/Oms
- CMakeLists.txt
- README.MD
>> mkdir build; cd build;
>> mkdir debug; cd debug; cmake -DCMAKE_BUILD_TYPE=Debug ../../; make -j4; ./Test/y; cd .. // compile and run test
>> mkdir release; cd release; cmake -DCMAKE_BUILD_TYPE=Release ../../; make -j4; ./Test/y // compile and run test
```

Compilation with Docker

First of all, `git clone` the *RHEL Docker* repository, which has three folders :

- all following `.in` files are seeds (i.e. initial files) to be processed into runnable scripts or config files
- most scripts take two arguments, which are the versions for *RHEL* and *gcc* respectively
- in order to use `RHEL7.2` image downloaded from *RHEL*, we have to register a developer account in *RedHat*

```
>> git clone git@ygit.yubo.local:developer/rhel-docker.git
>> cd rhel-docker; ll
- base/build          Python script for building RHEL7.2 + gcc10.2.0 image from RHEL7.2 image downloaded from RHEL
  base/Dockerfile.in  initial Docker file used by docker build
- local/build.sh      BASH script for building RHEL7.2 + gcc10.2.0 + YTL_user image from RHEL7.2 + gcc10.2.0 image
  local/Dockerfile    Docker file used by docker build
- project/configure    Python script for copying all utilities to YLibrary project folder
  project/bash.sh.in  utility to launch Docker container and run terminal
  project/build.sh.in utility to launch Docker container and run both cmake/make
  project/make.sh      utility to run cmake followed by make
  project/build/debug.sh final cmake command for debug (such as cmake -DCMAKE_BUILD_TYPE=Debug ../../)
  project/build/release.sh final cmake command for release (such as cmake -DCMAKE_BUILD_TYPE=Release ../../)
```

Secondly, build `RHEL7.2 + gcc10.2.0` image :

```
>> cd rhel-docker/base;
>> ./build 7.2 10.2.0
prompt for RHEL developer account
prompt for RHEL developer password
```

Thirdly, build `RHEL7.2 + gcc10.2.0 + YTL_user` image specific for *YLibrary*

```
>> cd rhel-docker/local;
>> ./build.sh 7.2 10.2.0
```

After that, copy utilities to *YLibrary* project folder by `configure` :

```
>> cd rhel-docker/project;
>> ./configure
```

After that, you can find a new `scripts` folder generated inside *YLibrary* project (they are massaged and copied from the seeds in *rhel-docker/project*), we can then build *YLibrary* using *RHEL Docker* container (simulate the environment in production) :

```
>> cd YLibrary/scripts; ll
- YLibrary/scripts/bash.sh
- YLibrary/scripts/build.sh
- YLibrary/scripts/make.sh
- YLibrary/scripts/build/debug.sh
  YLibrary/scripts/build/release.sh
>> ./build debug // cmake debug version and make it
>> ./build release // cmake release version and make it
>> ./build release y Test // cmake release version and make it, run Test exe
```

You can also run make in *nvim* by :

```
:set makeprg=scripts/build
:make debug
:make release
:make release y Test
```

Output will be shown in QuickFix.

Basic Docker

Docker concepts :

- Docker image - a file from which a container can be loaded
- Docker container - a box for running selected commands / tasks
- Docker registry - a Github for Docker image

```
// list all docker images, now we have 3 images
>> docker images
registry.redhat.io/rhel7:7.2    which is RHEL7.2 image
RHEL7.2:10.2.0                 which is RHEL7.2 + gcc10.2.0 image
YTL7.2:10.2.0                  which is RHEL7.2 + gcc10.2.0 + YTL_user image

// list all docker containers, now we have nothing, as we have not loaded any container yet
>> docker ps -a
nothing

// load docker container from image, then run desired command
>> docker run --rm --it registry.redhat.io/rhel7:7.2 bash
```

- option `--rm` means remove the container after completing the assigned task
- option `--it` means interactive mode, so that we can continue to use the terminal
- when Docker container is closed, all its files are gone, thus in order to save desired output ... we need to add mount points so that Docker containers can read / write files in local machine, using option `-v` and `-w`
- option `-v "$(pwd):$(pwd)"` means mounting current folder in local machine to current folder in container
- option `-w "$(pwd)"` means changing current directory to `$(pwd)`

```
>> docker run --rm -v "$(pwd):$(pwd)" -w "$(pwd)" ytl7.2:10.2.0 scripts/make.sh debug y Test
```

Basic scripting

<code>\$(abc)</code>	run command <code>abc</code>
<code>\$(pwd)</code>	run command <code>pwd</code> which returns current working directory
<code>\$(nproc)</code>	run command <code>nproc</code> which returns the number of processors
<code>\${@}</code>	all arguments of the BASH script
<code>\${0}</code>	the zeroth argument of the BASH script (i.e. BASH script name)
<code>\${1}</code>	the first argument of the BASH script
<code>\${2}</code>	the second argument of the BASH script
<code>\${#}</code>	number of arguments of the BASH script
<code>\${?}</code>	return code of the latest command
<code>\${abc}</code>	variable <code>abc</code>
<code>echo \${abc}</code>	print it
<code>\${BASH_SOURCE[0]}</code>	same as <code>\${0}</code> with minor differences, please check

Command BASH terminal command that are common in BASH scripting :

<code>>> shift</code>	pop away the first argument
<code>>> pushd path123</code>	equivalent to <code>cd path123</code> and push <code>pwd</code> into a directory stack
<code>>> popd</code>	equivalent to <code>cd directory_stack.front()</code> followed by pop
<code>>> dirname a/b/c</code>	return parent directory of input path, which is <code>a/b</code> in this case (<code>a/b/c</code> may not exist)

What's inside `YLibrary/scripts`?

There are 4/5 scripts in `scripts` folder. The first one `scripts/bash.sh` is an example of using *Docker* container. Here is its content :

```
#!/usr/bin/env bash
pushd "$(dirname "${BASH_SOURCE[0]}")/.." &> /dev/null
docker run --rm -it -v "$(pwd):$(pwd)" -w "$(pwd)" yt17.2:10.2.0 bash
popd &> /dev/null
```

- line 1 specifies `bash` interpreter
- line 2 jumps up two layers above `bash.sh`, which is folder `YLibrary`
- line 3 loads *Docker* container, create mount point to folder `YLibrary` in local machine from exactly the same point in container
- line 4 goes back to `bash.sh`

The second one `scripts/build.sh` is similar to `scripts/bash.sh`, which is specific to `cmake` and `make` only :

```
#!/usr/bin/env bash
pushd "$(dirname "${BASH_SOURCE[0]}")/.." &> /dev/null
docker run --rm -v "$(pwd):$(pwd)" -w "$(pwd)" yt17.2:10.2.0 scripts/make.sh ${@}
rc=${?}
popd &> /dev/null
exit ${rc}
```

- line 4 and 6 are extra (compared to `bash.sh`), they get the return code from `scripts/make.sh` and return
- line 3 includes `${@}`, which means forwarding all arguments of `build.sh` to `make.sh`
- in other words, `build.sh` is run in local machine, while `make.sh` is run in container

The third one `scripts/make.sh` is the one who invokes (1)`cmake` and (2)`make` :

```
#!/usr/bin/env bash
if [[ $# -lt 1 ]]; then
    exit 1
fi

BUILD_NAME="${1}"
BUILD_DIR="build/${1}"
shift

pushd "$(dirname "${BASH_SOURCE[0]}")/.." &> /dev/null
mkdir -p "${BUILD_DIR}"
pushd "${BUILD_DIR}" &> /dev/null

# step 1
ulimit -c unlimited
CMAKE="../../scripts/build/${BUILD_NAME}.sh"
cat "${CMAKE}" | grep cmake
"${CMAKE}"
rc=${?}
if [[ $rc -ne 0 ]]; then
    exit $rc
fi

# step 2
make -j ${nproc} ${@}
rc=${?}
if [[ $rc -ne 0 ]]; then
    exit $rc
fi

popd &> /dev/null
popd &> /dev/null
```

-lt means less-than
red characters denote the IF-CONDITION

BUILD_NAME = debug, release, production
BUILD_DIR = build/debug, build/release, build/production
pop away : debug, release, production

set coredump limit for Test

display the command to terminal or nvim's quickfix
execute the command
return code of the command
return code equals to zero for success, quits if cmake fails

Finally, `scripts/build/*.sh` is just `cmake` invocation :

```
#!/usr/bin/env bash
cmake -DCMAKE_C_COMPILER=$(which gcc)
      -DCMAKE_CXX_COMPILER=$(which g++)
      -DCMAKE_BUILD_TYPE=Debug ../../
```

KAFKA

For topology of KAFKA, please read <http://cloudurable.com/blog/kafka-architecture/index.html>

How to install and run KAFKA?

1. visit <https://kafka.apache.org/quickstart> and download KAFKA
2. unzip and install as follows ...

```
>> cd ~/Downloads
>> tar -xzf kafka_2.13-2.8.0.tgz
```

Zookeeper is for :

- coordination, management of KAFKA servers (clusters)
- including load balancing
- including recovery

KAFKA is for :

- KAFKA server = cluster
- KAFKA client = producer / client
- producer = publisher
- consumer = subscriber

Lets view and modify config :

```
>> cd kafka_2.13-2.8.0
>> nvim config/zookeeper.properties (nothing changed)
>> nvim config/server.properties
```

Add two lines at the end of server.properties :

```
auto.create.topics.enable = false
delete.topic.enable = true
```

Steps in sequence in different terminal (you can use `tmux`) :

1. start zookeeper in terminal 0
2. start KAFKA server in terminal 1
3. create KAFKA topics in terminal 2 (terminal 2 can be re-used)
4. start client console – producer in terminal 2
5. start client console – consumer in terminal 3 (both producer and consumer are KAFKA clients)

All the executables (bash scripts) are available here :

```
>> ll bin
```

Step 1&2 : Run zookeeper and KAFKA server with appropriate config and options :

```
>> bin/zookeeper-server-start.sh config/zookeeper.properties
>> bin/kafka-server-start.sh config/server.properties
```

Step 3 : Create multiple topics

```
>> bin/kafka-topics.sh --bootstrap-server localhost:9092 \      (identifier of KAFKA server)
    --create \
    --topic my_topic_ABC \
    --replication-factor 1 \
    --partitions 1
>> bin/kafka-topics.sh --bootstrap-server localhost:9092 \      (identifier of KAFKA server)
    --list \                                                    (list all topics)
>> bin/kafka-topics.sh --bootstrap-server localhost:9092 \      (identifier of KAFKA server)
    --describe --topic my_topic_ABC \                            (list one topic)
```

Step 4&5 : Start two client instances, one for producer and one for client :

```
>> bin/kafka-console-producer.sh --bootstrap-server localhost:9092 \
    --topic my_topic_ABC
>> bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 \
    --topic my_topic_ABC
>> bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 \
    --topic my_topic_ABC \
    --group group_name_DEF
```

where :

```
--bootstrap-server    specifies ip:port of KAFKA server (port 9092 can be found in config)
--create              option to create topic
--list                option to list topic
```

Stop clients by ctrl-c and stop servers by :

```
>> bin/kafka-server-stop.sh
>> bin/zookeeper-server-stop.sh
```

KAFKA vs database

About database concepts :

ACID = xxx-ability : atomicity / consistency / isolation / durability

CRUD = xxx-operation : create / read / update / delete

CRUD	SQL	HTTP	STL
create	insert	post	insert
read	select	get	operator[] const
update	update	put	operator[]
delete	delete	delete	erase

ELK stack =

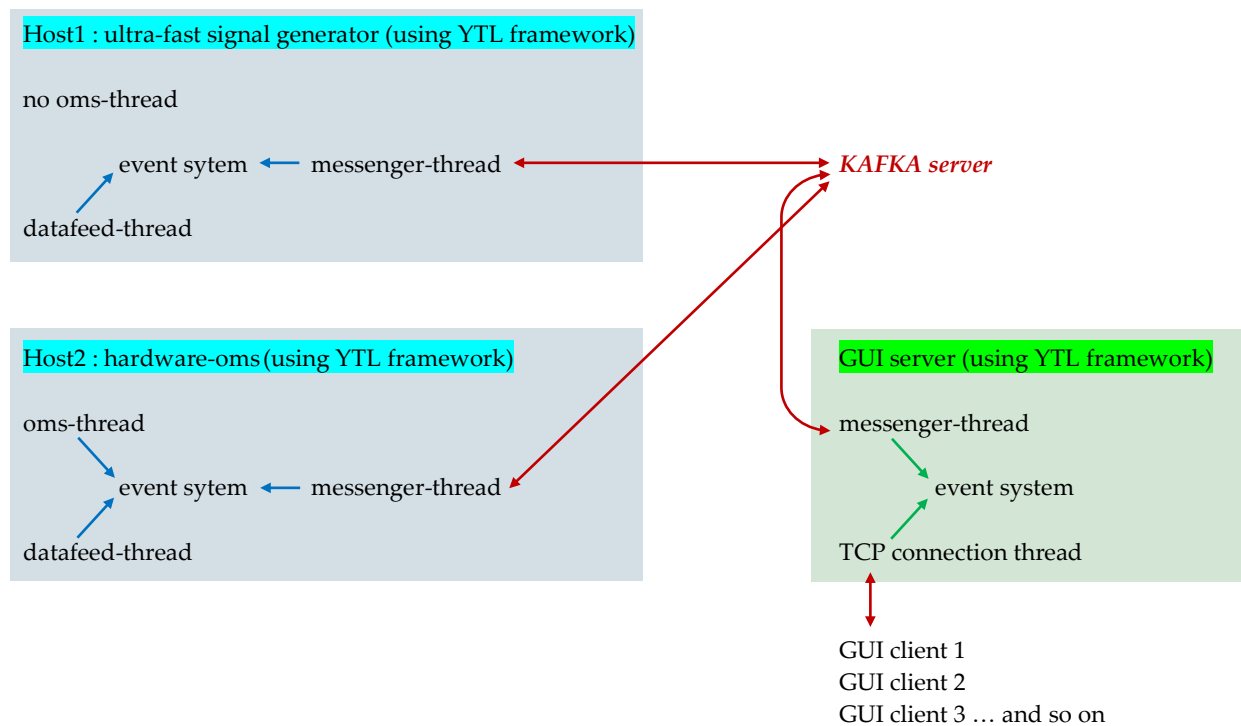
Elastic search +

Log stash +

Kibana

Here is the diagram showing integration of :

- event system (inter-threads in same machine) with
- KAFKA pub-sub sytem (inter-machines)



In the above architecture, there are two hosts :

- signal generator can estimate theoretical-value of underlying in ultra-fast way making use of TCP checksum to price mapping
- hardware-oms can emit buy and sell order when market price is more favorable to theoretical-value from signal generator

Please trace the route :

- from GUI client to GUI server to KAFKA and finally to host
- from host to KAFKA to GUI server and finally back to GUI client
- on both directions, it involves several event systems